

Transformers and Attention

How attention works, why it replaced recurrence, and what actually happens inside a large language model, without the mystique.



Utsav Poudel
Founder, Fulcrum

ABOUT FULCRUM

Who made this, and why it is free

Fulcrum is a volunteer-run non-profit. We help people from under-resourced regions learn artificial intelligence through one-on-one mentorship, research and publication supervision, project support, workshops, and a guest speaker series.

The premise is simple. Ability is spread evenly across the world; opportunity is not. What separates a capable student in Pokhara or rural Bihar from one at a well-funded lab is rarely intelligence. It is access to supervision, to reviewers, to equipment, and to the unwritten rules nobody writes down.

This deck is part of how we try to close that. It is free, it will stay free, and you may reuse and adapt it for your own teaching under CC BY 4.0. If it is useful to you, come and find us.

Find us

LinkedIn

linkedin.com/company/gofulcrum

Discord

discord.gg/gbQCGkupdJ

Instagram

[@gofulcrum](https://www.instagram.com/gofulcrum)

Every program is free. Mentorship applications go through Discord.

PART ONE

The problem attention was invented to solve

Recurrence forces information through a single hidden state, one step at a time. Both halves of that are costly.

What was wrong with recurrence

- ▲ An RNN compresses everything seen so far into one fixed-size hidden state, a bottleneck for long inputs.
- ▲ It is **inherently sequential**, so it cannot exploit parallel hardware during training.
- ▲ Long-range dependencies degrade: gradients must travel through every intervening step.
- ▲ LSTMs and GRUs mitigated this with gating, but did not remove either limitation.

Attention, stated plainly

Let every position look directly at every other position, and learn how much to weight each.

- ▲ Each token produces three vectors: a **query**, a **key**, and a **value**.
- ▲ Compare one token's query against all keys to get a relevance score per token.
- ▲ Softmax those scores into weights; the output is the weighted sum of values.
- ▲ Nothing is compressed into a bottleneck, and every comparison happens in parallel.

Resolving a pronoun

every token looks at every other token



Resolving "it" means weighting "cat" heavily and the rest lightly. Those weights are learned.

The formula, decoded

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d}) V$$

- ▲ QK^T , every query dotted with every key, giving a score matrix of shape $n \times n$.
- ▲ $/ \sqrt{d}$, scaling by the square root of head dimension. Without it, large dot products push softmax into saturation and gradients vanish.
- ▲ softmax , turns scores into weights that sum to one.
- ▲ V , the weighted sum of values is the output for that position.

That $n \times n$ matrix is the whole efficiency story: cost grows with the square of sequence length. Everything called 'efficient attention' is an attack on that term.

Multi-head attention

- ▲ Run several attention operations in parallel, each with its own learned projections.
- ▲ Different heads specialise. Some track syntax, some track coreference, some attend to position.
- ▲ Concatenate the heads' outputs and project back to the model dimension.
- ▲ Each head works in a smaller subspace, so total cost is comparable to one full-width head.

Head interpretations are suggestive, not definitive. Be careful reading too much meaning into attention maps. Attention weight is not the same thing as causal importance.

The rest of the block

- ▲ **Positional encoding**, attention is permutation-invariant, so position must be injected explicitly. Sinusoidal originally; rotary (RoPE) is now common.
- ▲ **Feed-forward network**, a per-position MLP, typically $4\times$ the model width. This is where most parameters live.
- ▲ **Residual connections** around both sub-layers, plus **layer normalisation**.
- ▲ Stack this block N times. That is the entire architecture.

PART TWO

From architecture to language model

The architecture is simple. The behaviour comes from scale and training objective.

Encoder, decoder, or both

- ▲ **Encoder-only** (BERT), bidirectional context, trained by masking tokens. Good for classification and retrieval.
- ▲ **Decoder-only** (GPT family), causal masking so each position sees only the past. Trained to predict the next token. Good for generation.
- ▲ **Encoder-decoder** (T5, original Transformer), natural fit for translation and summarisation.
- ▲ Decoder-only has become dominant largely because next-token prediction scales so cleanly.

Tokenisation

- ▲ Models operate on **subword tokens**, not characters or words, via byte-pair encoding or similar.
- ▲ This keeps vocabulary manageable while handling unseen words by composition.
- ▲ Tokenisers are trained mostly on English-dominant corpora, so other scripts often fragment into far more tokens.
- ▲ That is a real equity issue: the same sentence in Nepali can cost several times more tokens than in English, making it slower and more expensive to process.

Worth checking yourself

Run a sentence in your own language through a tokeniser and count. It is a concrete, publishable observation about who these systems are cheap for.

How training actually proceeds

- ▲ **Pretraining**, next-token prediction over an enormous corpus. This is where almost all compute goes.
- ▲ **Supervised fine-tuning**, training on curated instruction-and-response pairs to make the model follow requests.
- ▲ **Preference tuning** (RLHF, DPO), optimising against human preference comparisons.
- ▲ The base model has the knowledge; the later stages shape how it behaves.

What these models do and do not do

- ▲ They model the **distribution of text**. Fluency is the objective; truth is not.
- ▲ **Hallucination** is not a bug bolted on top. It is the same mechanism producing a plausible continuation that happens to be false.
- ▲ **Context window** bounds what the model can attend to. Beyond it, information is simply gone.
- ▲ They are strong at transformation (summarising, rewriting, translating) and much weaker as a factual store.

Useful framing for teaching: it is an extremely good pattern completer. That explains both the impressive results and the confident errors.

Working with these models on modest hardware

- ▲ **Quantisation** to 4-bit or 8-bit lets sizeable open models run on a single consumer GPU.
- ▲ **LoRA** trains small low-rank updates instead of full weights, fine-tuning becomes affordable.
- ▲ **Retrieval-augmented generation** grounds answers in documents you supply, which is usually better than fine-tuning facts in.
- ▲ **Prompting well** costs nothing and is frequently the highest-return intervention available.

Open questions worth a paper

- ▲ Tokenisation and cost equity across scripts and low-resource languages.
- ▲ Evaluation beyond English. Most benchmarks do not transfer, and building one for your language is a genuine contribution.
- ▲ Small-model performance under tight compute, which matters far more outside well-funded labs.
- ▲ Careful failure analysis of deployed systems in specific domains. Unglamorous, useful, and publishable.

Choosing a question

The frontier is crowded and expensive. The gaps are in places large labs have no reason to look, which is precisely where our mentees live.

ABOUT FULCRUM

Work on one of these with us

Fulcrum mentors supervise research from question to submission. Free, and open to anyone from an under-resourced region.

Fulcrum is a volunteer-run non-profit helping people from under-resourced regions learn AI through free mentorship, research supervision, and open material like this deck. Founded 2025 in Kathmandu, Nepal. Nothing we do costs anything.

Apply for mentorship

discord.gg/gbQCGkupdJ

Partnerships and enquiries

linkedin.com/company/gofulcrum

Instagram

[@gofulcrum](https://www.instagram.com/gofulcrum)